


Meet the companion who reads your sky out loud.

- ✧ A natal astrology and synastry companion.
- ✧ Powered by on-device inference and bespoke local memory.
- ✧ Designed for spoken conversation, not endless scrolling.



Breaking the boundary between calculation and connection.



Traditional apps treat astrology as a static text delivery system. Standard chatbots treat interaction as a transactional search. Natally treats the user as a patron sitting across the table from a friend.

The fortune-teller persona is the primary interface.

Direct Eye Contact

Fosters emotional intimacy and active listening.

The Crystal Ball

Serves as the dynamic focal point where chart generation and 'thinking' states are visualized.

Warm Kintsugi Elements

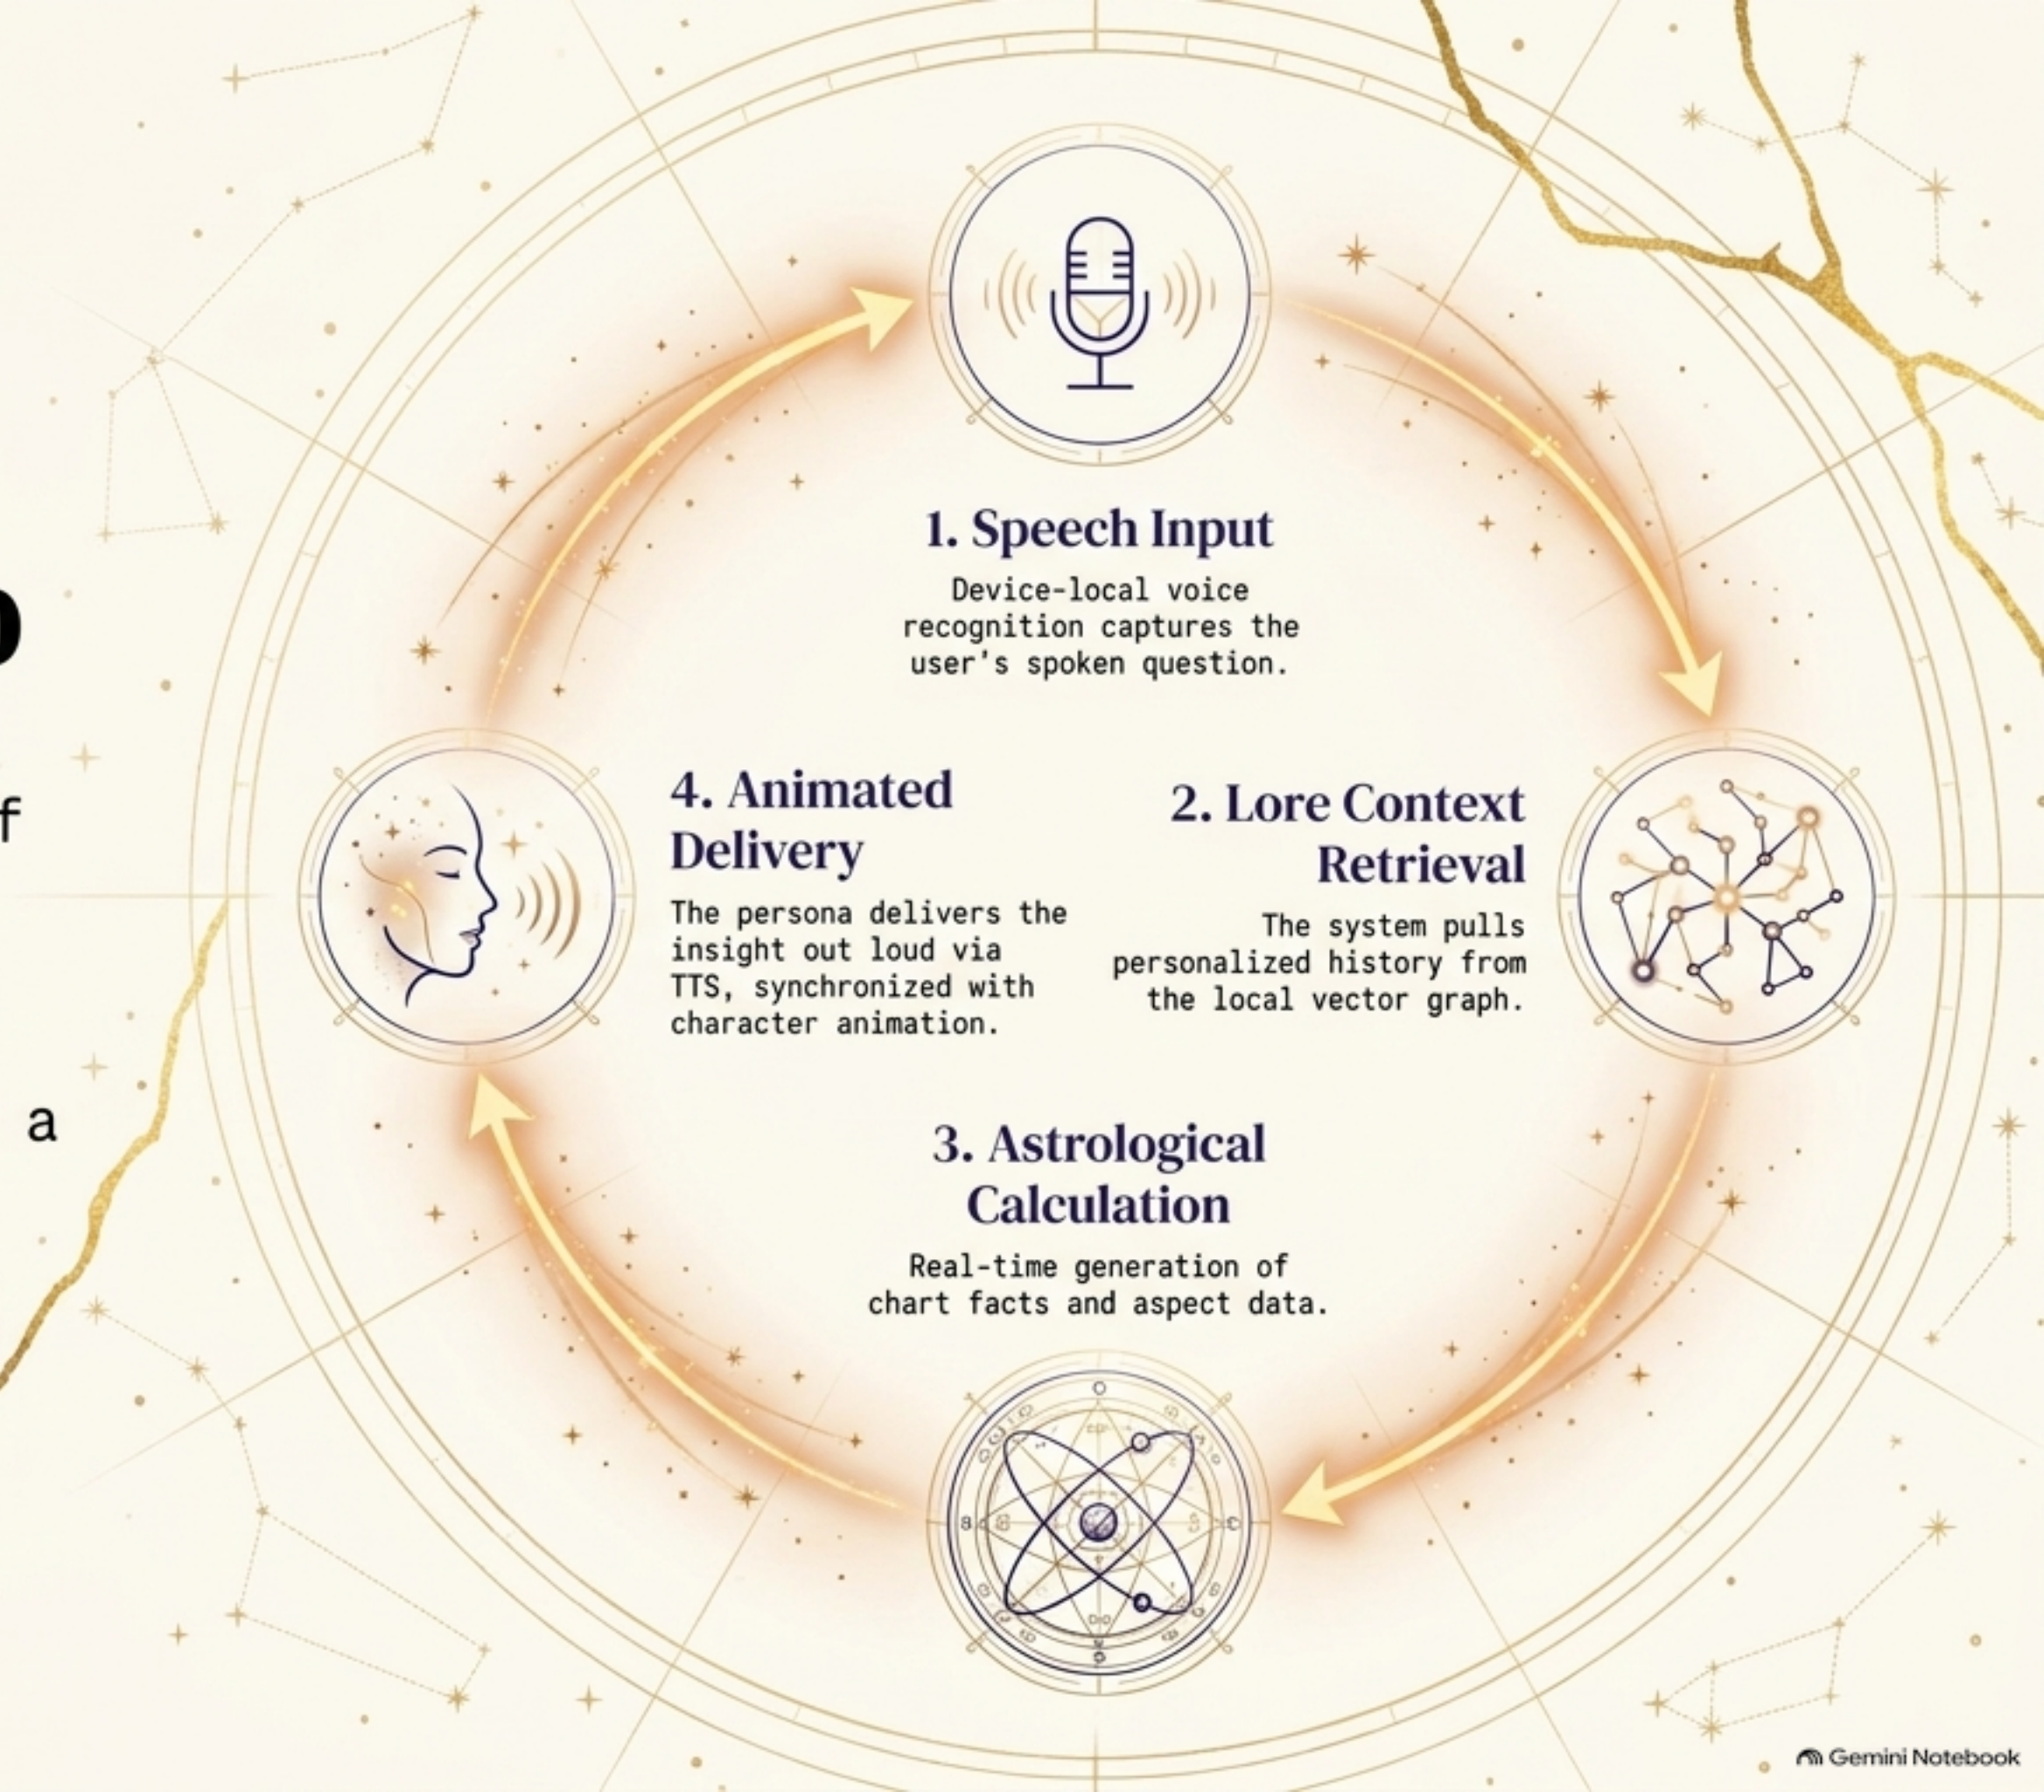
Golden seams in the design language visually reinforce the theme of healing, memory, and personalized care.



Natally is a person, not a dashboard. The app features no complex nested menus to navigate. Charts and planetary transits are simply the things she shows you while she talks to you.

A continuous loop of spoken companionship

Typing breaks the illusion of presence. Natally operates through continuous spoken interaction, transforming rigorous data retrieval into a natural, listening relationship.



Immersive design rooted in visual storytelling.

Every screen leverages shared UI source code (Figma-to-Stitch HTML/CSS/JS), ensuring the bold type, intricate motion motion design, and warm conversational feel persist flawlessly across both local devices and hosted web environments.

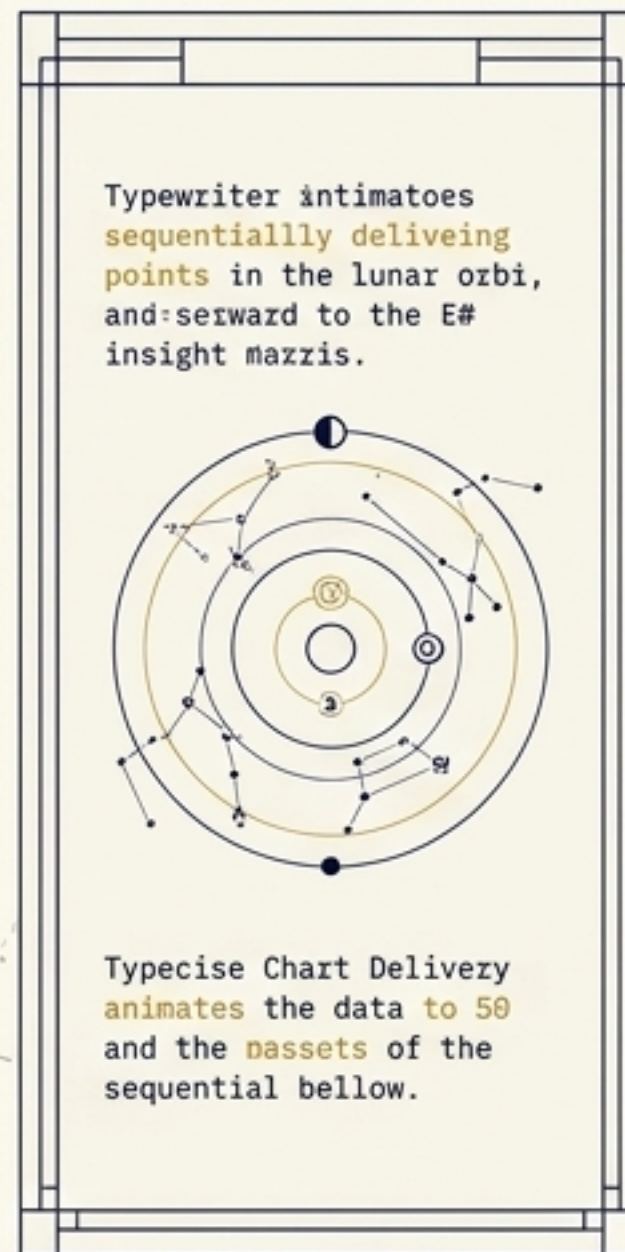
The Kintsugi Onboarding



Thinking Pysanky Background

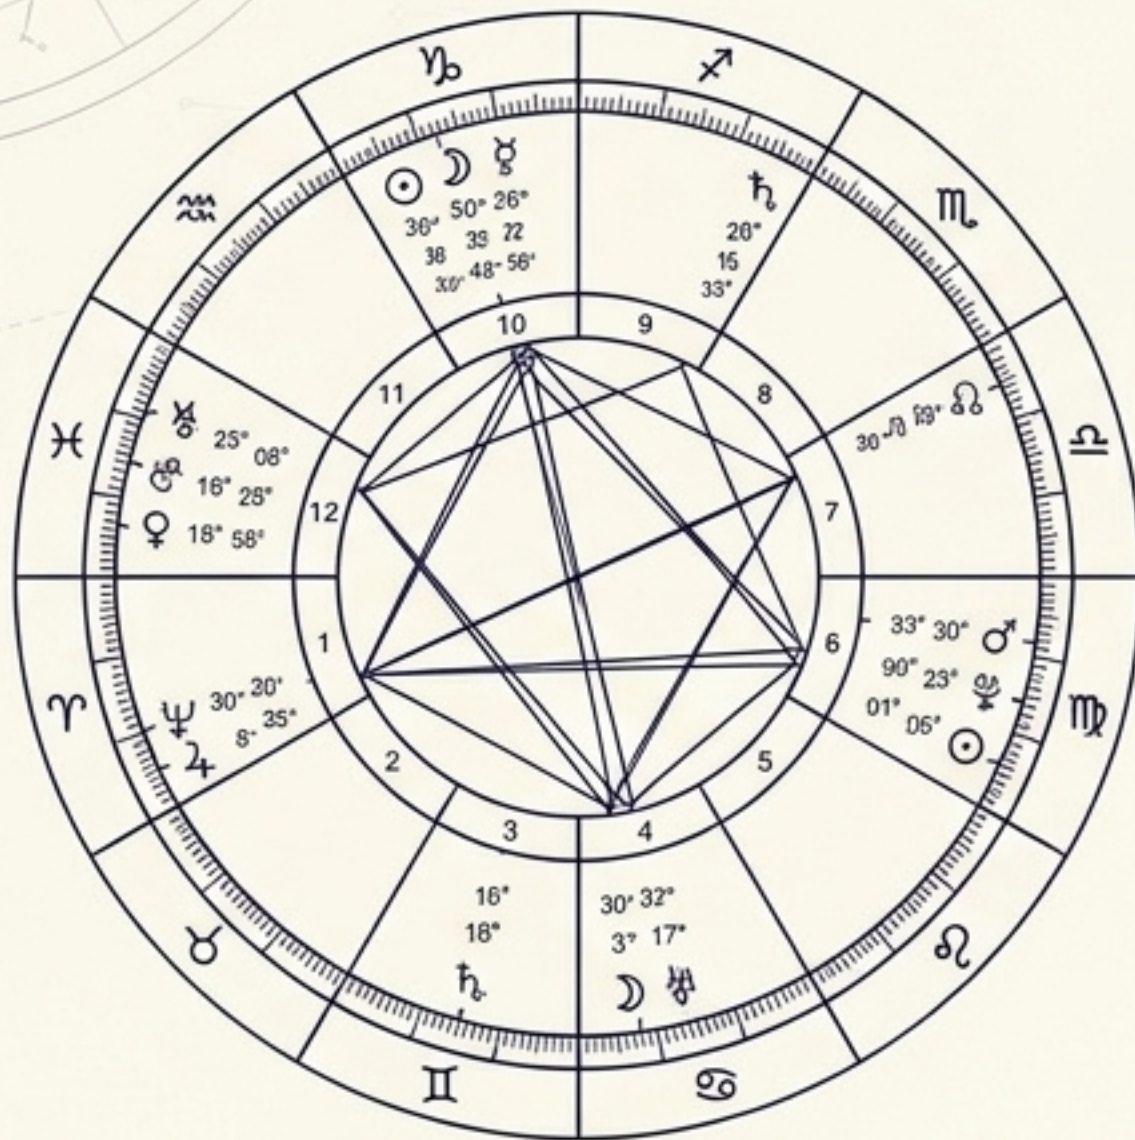


Precise Chart Delivery

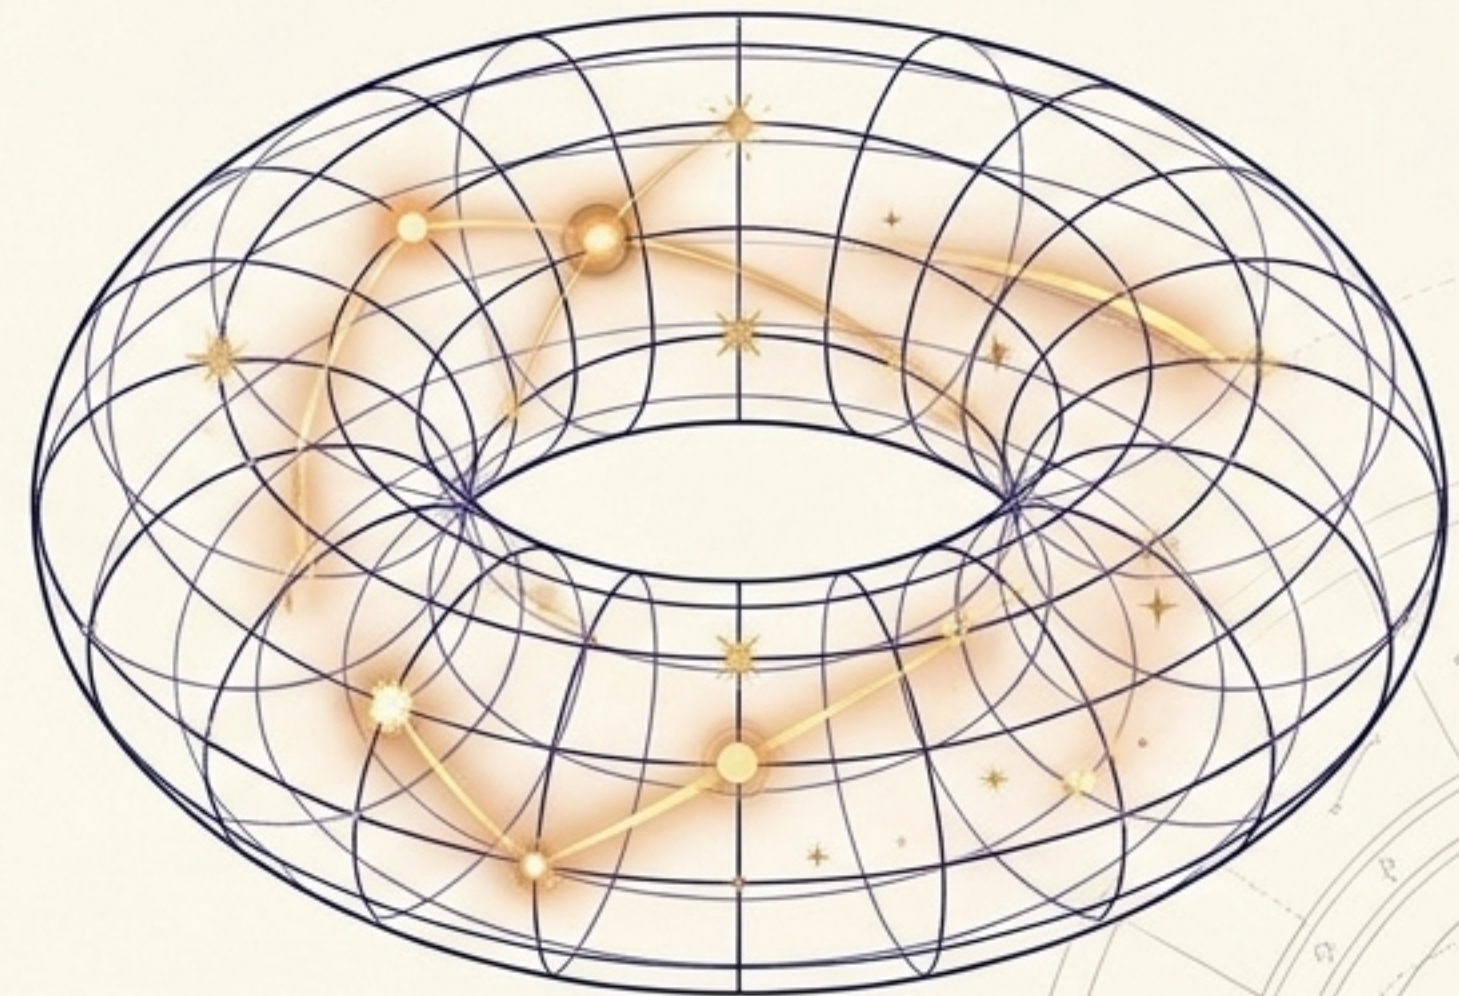


Visualizing the sky with mathematical precision.

Behind the intuitive conversation lies a rigorous astrological engine. Whether calculating personal natal placements, synastry (relationship) compatibility, or daily transits, the companion projects the exact mathematics of the sky into stunning 3D visualizations.



Standard 2D Cartography



3D Torus Projection

Contextual memory powered by local vector embeddings.

Companionship requires shared history. Every conversation, extracted chart fact, and stated preference is embedded into a local knowledge graph. This "lore" acts as the brain behind the persona, making every session deeply and continuously personalized.



One monorepo. Two parallel deployments.

Natally's engineering triumph is its parallel build structure. The unified monorepo ensures the exact same rich, companion-led design and tool-calling boundaries are deployed flawlessly whether running entirely offline on a phone or via hosted cloud architecture.

apps/local

- Device Inference
- Native runtime
- Offline-capable PWA

apps/hosted

- Web App
- Hosted Cloud Inference
- API Services

Shared Figma/Stitch Exported UI
Source & Companion Experience

Architecture comparison: The Ecosystem.

Two environments, one identical user experience. The local edition prioritizes absolute privacy and offline autonomy, while the hosted edition guarantees immediate, zero-download access.

	Local Edition	Hosted Edition
Target Environment	Android, Linux, Windows, browser PWA	Web browser via URL
Connectivity	Fully offline capable after initial asset installation	Requires continuous internet connection
Data Privacy	Lore and transcripts stored entirely on-device	Client-side lore context sent securely to API
Execution Runtime	Native llama.cpp / wllama WASM	Cloud-hosted server runtime

Architecture comparison: The Intelligence.

Running a 2B parameter model locally establishes the floor for true companionship—handling complex tool calling and contextual lore decisions natively—while the hosted edition scales intelligence through connected cloud APIs.

	Local Edition	Hosted Edition
Core Inference (LLM)	Qwen3.5-2B-Q4_K_M.gguf (DOM access + tool calling capabilities)	OpenRouter (dynamic backoff, cloud-tier weights)
Voice Generation (TTS)	Kokoro-82M v1 (ONNX synthesis, native playback)	Premium cloud speech generation API
Lore Embeddings	all-minilm-16-v2-q8_0	Shared browser-local embeddings wired to API

The strategic evolution of the ephemeris engine.

Phase 1 (Current): Open Source

- AGPL-3.0-or-later
- Utilizing sweph-wasm (Swiss Ephemeris)
- Isolated within the VENDORED/ directory

Phase 2 (Target): Proprietary

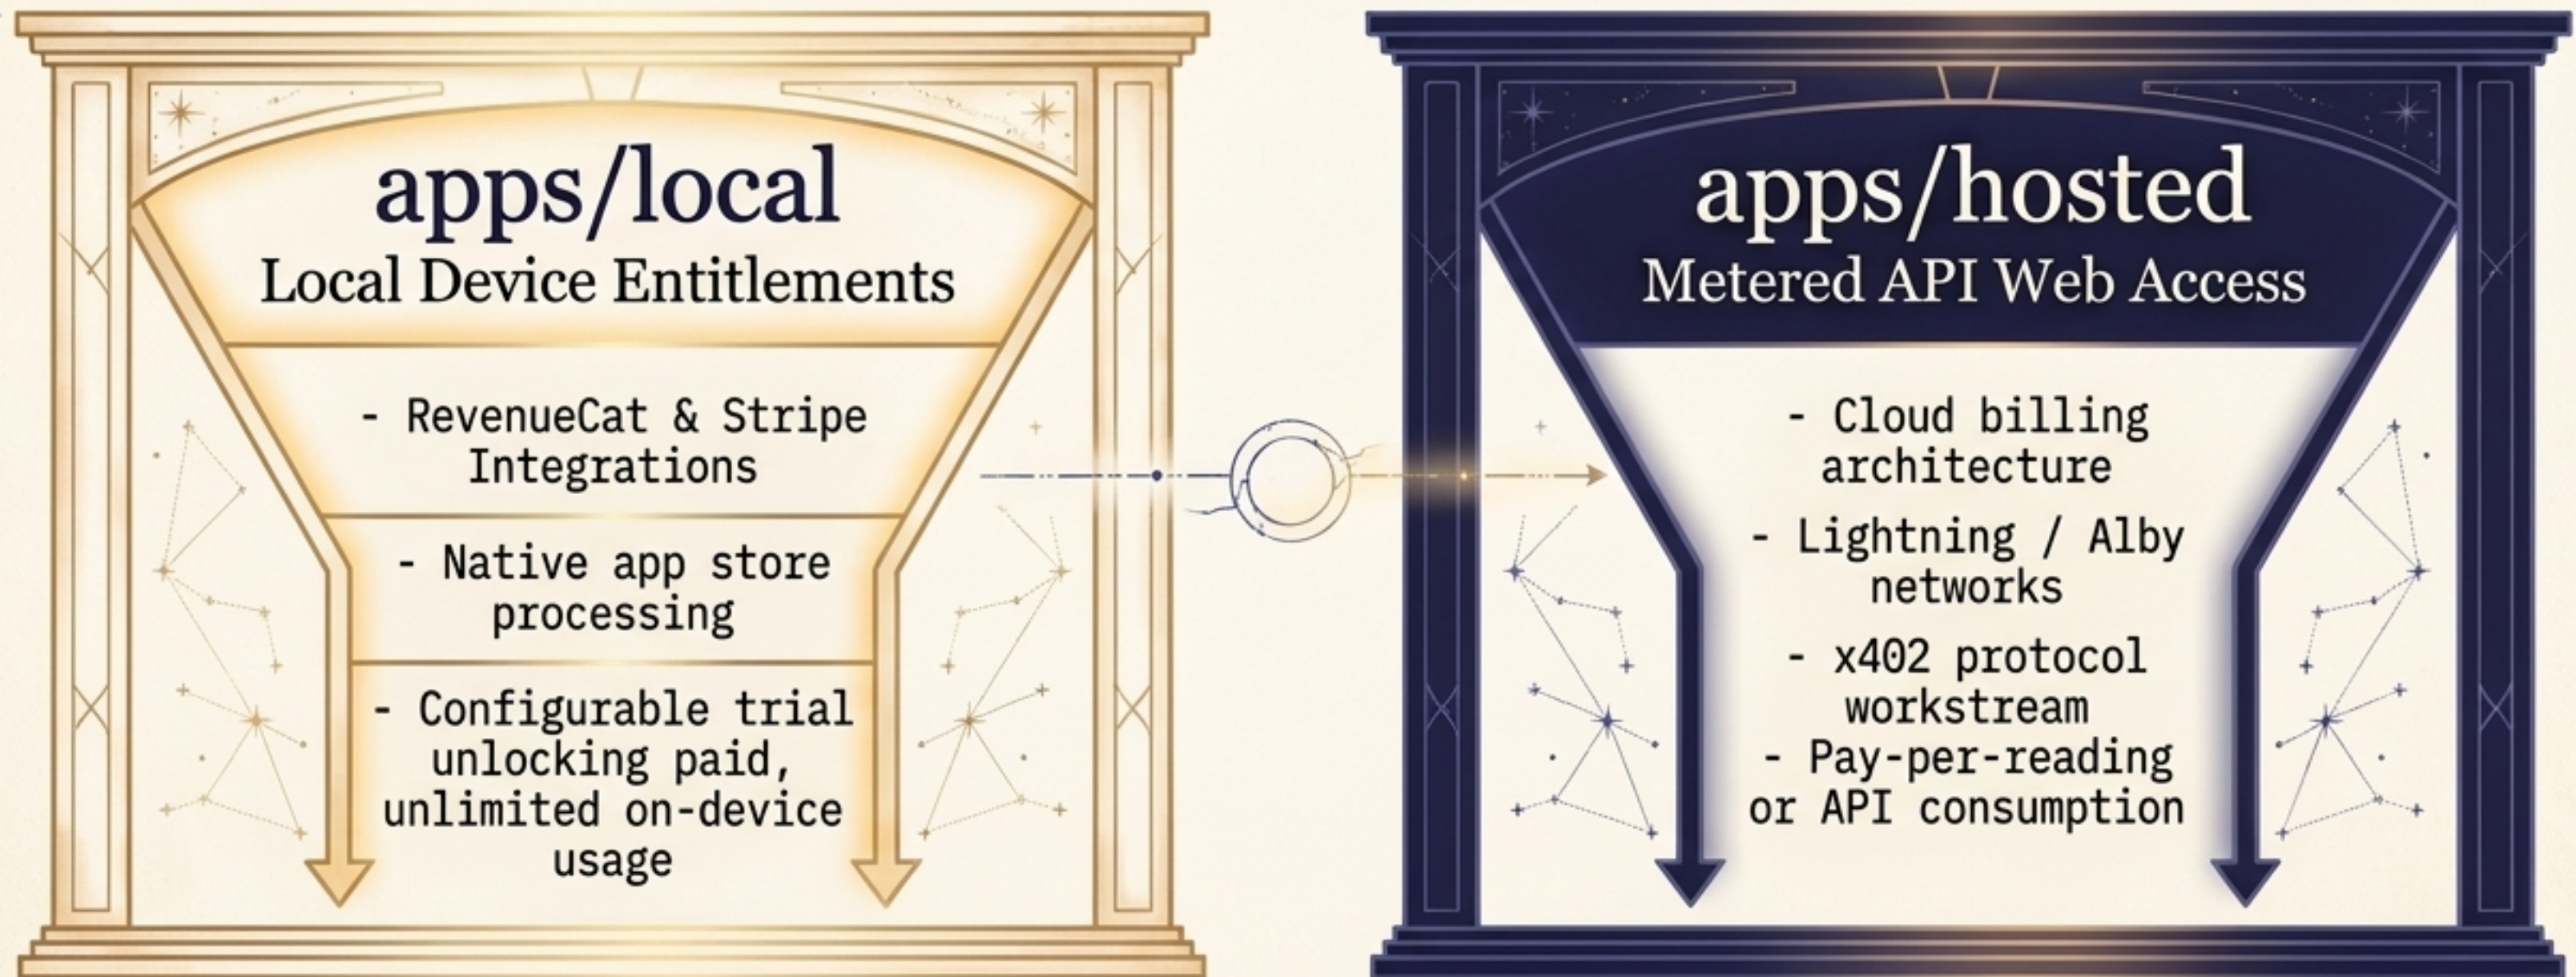
- Bespoke in-house ephemeris engine supersedes the Swiss Ephemeris
- Allows for commercial licensing flexibility

The Seam:
EphemerisEngine
Contract Boundary

To ensure long-term commercial viability, the architecture isolates the current AGPL-licensed Swiss Ephemeris behind a strict interface contract. This cleanly paves the way for a seamless drop-in replacement with a proprietary, in-house mathematical engine.

Flexible access models for parallel ecosystems.

The deployment strategy inherently de-risks the business model. By maintaining local entitlements alongside metered cloud access, Nattally captures both the privacy-conscious enthusiast willing to pay for software ownership, and the casual web user requiring instant, frictionless entry.



Where astrophysics and local engineering disappear behind the face of a friend.



Natally represents a paradigm shift in human-computer interaction. It proves that with an uncompromising monorepo architecture, optimized local-LLM inference, and rich visual design, the most complex technical systems can feel exactly like a conversation with a companion.